TELECOM CUSTOMER CHURN ANALYSIS

Predictive Modelling and Business Intelligence Report

Dataset: IBM Telco Customer Churn
Observations: 7,043 | Features: 21 | Target: Churn (Yes/No)

Methods: EDA | Statistical Testing | Logistic Regression | Random Forest | XGBoost | SHAP

March 2026

1. Executive Summary

This report presents a comprehensive data-driven analysis of customer churn in a telecom dataset comprising 7,043 customers and 21 variables. The analysis proceeds through five methodological stages: exploratory data analysis (EDA), formal statistical testing, interpretable logistic regression modelling, ensemble methods (Random Forest and XGBoost), and model explainability via SHAP values.

The overall churn rate is 26.5%, exceeding the industry average of 15-20% and indicating a significant retention problem. Across all analytical methods, four variables consistently emerge as the primary churn drivers: contract type, customer tenure, internet service type, and monthly charges. These findings are statistically robust, cross-validated, and methodologically consistent.

The best-performing models are Logistic Regression (AUC = 0.845 +/- 0.014) and XGBoost (AUC = 0.847 +/- 0.021), both achieving Recall approximately 0.79 under 10-fold cross-validation. Five evidence-based recommendations are provided for immediate business implementation.

Churn Rate	Best AUC (CV)	Best Recall (CV)	Features Used
26.5%	**0.847 +/- 0.021**	**0.792 +/- 0.035**	**15 of 19**

2. Dataset & Preprocessing

2.1 Dataset Overview

The dataset contains 7,043 customer records across 21 variables including demographic information, service subscriptions, billing details, and churn status. No missing values were present in the original dataset. The target variable Churn is binary (Yes/No), with a class distribution of 73.5% No Churn and 26.5% Churn.

2.2 Variable Exclusions

Four variables were excluded prior to modelling based on statistical and theoretical grounds:

Variable	Reason	Evidence
TotalCharges	Multicollinearity	r = 0.826 with tenure (exceeds threshold of 0.80)
customerID	Non-predictive	Unique identifier with no predictive value
gender	Non-significant	Chi-square p = 0.487 (fail to reject H0)
PhoneService	Non-significant	Chi-square p = 0.339 (fail to reject H0)

2.3 Encoding & Scaling

- **Label Encoding:** Binary variables (Yes/No): Label Encoding (0/1)
- **One-Hot Encoding:** Variables with 3+ categories: One-Hot Encoding with drop_first=True
- **StandardScaler:** Applied StandardScaler to tenure and MonthlyCharges
- **Data Leakage Prevention:** Fit on training set only; applied to test set to prevent data leakage

2.4 Train/Test Split

- 80% training (5,634 rows) / 20% test (1,409 rows)
- Stratified split to maintain 26.5% churn ratio in both sets
- random_state = 42 for full reproducibility

3. Exploratory Data Analysis

3.1 Churn Distribution

The dataset exhibits a 73.5% / 26.5% class split (5,174 No Churn vs 1,869 Churn). This constitutes a moderately imbalanced dataset (approximately 3:1 ratio), necessitating the use of class_weight='balanced' in all classification models. The churn rate of 26.5% substantially exceeds the telecom industry average of 15-20%, indicating a systemic retention problem.

3.2 Churn vs Tenure

Median tenure for churned customers is 10 months versus 38 months for retained customers. The histogram reveals a sharp spike in churn during months 1-2, identifying a critical early lifecycle vulnerability window. Customers who survive beyond 36 months demonstrate near-zero churn probability. This pattern was confirmed statistically by the Mann-Whitney U test ($U = 7{,}154{,}668$, $p < 0.001$).

3.3 Churn vs Contract Type

Contract Type	Churn Rate	Business Implication
Month-to-month	**42.7%**	Highest risk — no contractual barrier
One year	11.3%	Moderate risk — annual commitment helps
Two year	**2.8%**	Lowest risk — strong retention effect

3.4 Churn vs Monthly Charges

Churned customers pay a median of $79.65 per month compared to $64.43 for retained customers (Mann-Whitney $U = 3{,}667{,}080$, $p < 0.001$). The violin plot reveals a bimodal distribution among retained customers: one cluster at approximately $20 (basic plans) and one at approximately $85 (satisfied premium customers). Churned customers are concentrated in the $60-$100 range, suggesting a perceived value-price misalignment at mid-to-high price points.

4. Statistical Testing

4.1 Normality Testing

Shapiro-Wilk tests confirmed non-normality for all numerical variables in both churn groups ($p < 0.001$ in all cases). Consequently, the non-parametric Mann-Whitney U test was selected for numerical variable comparisons, rather than the parametric t-test.

4.2 Numerical Variables — Mann-Whitney U Test

Variable	Median (No Churn)	Median (Churn)	p-value	Significant
tenure	38.00	10.00	< 0.001	Yes
MonthlyCharges	64.43	79.65	< 0.001	Yes

4.3 Categorical Variables — Chi-Square Test

Chi-square tests of independence were conducted for all 16 categorical variables. Expected cell frequencies exceeded 5 in all cases, satisfying the test assumption. Results are summarised below for key variables:

Variable	Chi-Square	df	p-value	Significant
Contract	1184.60	2	< 0.001	Yes
OnlineSecurity	850.00	2	< 0.001	Yes
TechSupport	828.20	2	< 0.001	Yes
InternetService	732.31	2	< 0.001	Yes
gender	0.48	1	0.487	No
PhoneService	0.92	1	0.339	No

17 of 19 variables were statistically significant at alpha = 0.05. The two non-significant variables (gender, PhoneService) were excluded from all predictive models.

5. Logistic Regression

5.1 Model Specification

- Algorithm: Logistic Regression (sklearn + statsmodels for inference)
- Class weight: balanced (accounts for 73/27 imbalance)
- Solver: lbfgs (sklearn), bfgs (statsmodels)
- Pseudo R-squared (McFadden): 0.2768 — indicating good model fit
- LLR p-value: < 0.001 — model significantly better than null

5.2 Key Coefficients with 95% Confidence Intervals

The following table presents statistically significant coefficients ($p < 0.05$) from the statsmodels implementation, which provides formal p-values and confidence intervals via maximum likelihood estimation. Odds Ratios (OR) are computed as exp(coefficient) and represent multiplicative changes in churn odds.

Feature	Coeff.	OR	CI Lower	CI Upper	p
Contract_Two year	-1.524	0.218	0.150	0.317	<.001
tenure	-0.802	0.449	0.395	0.510	<.001
Contract_One year	-0.796	0.451	0.357	0.570	<.001
OnlineSecurity_Yes	-0.368	0.692	0.565	0.847	<.001
TechSupport_Yes	-0.327	0.721	0.588	0.884	.002
InternetService_Fiber optic	+0.561	1.753	1.189	2.584	.005
PaymentMethod_Electronic check	+0.437	1.548	1.259	1.904	<.001
PaperlessBilling	+0.427	1.533	1.305	1.802	<.001

Green rows indicate churn-reducing factors (OR < 1); red rows indicate churn-increasing factors (OR > 1). All 95% confidence intervals exclude 1.0, confirming statistical significance.

5.3 Non-Significant Features (p >= 0.05)

Six features failed to achieve statistical significance in the full model: MultipleLines_Yes ($p = 0.052$), PaymentMethod_Mailed_check ($p = 0.851$), Partner ($p = 0.813$), DeviceProtection_Yes ($p = 0.979$), PaymentMethod_Credit_card_automatic ($p = 0.875$), and OnlineBackup_Yes ($p = 0.240$). These variables do not contribute independent predictive value when other variables are controlled for.

6. Model Comparison

6.1 Models Evaluated

- Logistic Regression — interpretable baseline with formal statistical inference
- Random Forest — ensemble of 200 decision trees trained in parallel
- XGBoost — sequential gradient boosting with 200 trees (learning rate = 0.05)

All models used class_weight='balanced' or equivalent (scale_pos_weight) to address class imbalance. All were evaluated on the identical test set (n = 1,409) and via 10-fold stratified cross-validation on the training set (n = 5,634).

6.2 Test Set Performance

Model	Accuracy	Precision	Recall	F1-Score	ROC-AUC
Logistic Regression	0.743	0.511	**0.778**	0.617	0.839
Random Forest	**0.781**	**0.613**	**0.479**	**0.538**	0.820
XGBoost	0.746	0.514	**0.783**	**0.621**	**0.842**

Green cells indicate the best value per column; red cells indicate the worst. Random Forest achieves the highest accuracy and precision but critically fails on Recall (0.479), missing 52% of actual churners — an unacceptable result for a churn detection use case.

6.3 10-Fold Cross-Validation Results

Model	AUC (mean +/- std)	Recall (mean +/- std)	F1 (mean +/- std)
Logistic Regression	0.845 +/- 0.014	**0.792 +/- 0.035**	0.626 +/- 0.025
Random Forest	**0.818 +/- 0.016**	**0.477 +/- 0.038**	**0.541 +/- 0.029**
XGBoost	**0.847 +/- 0.021**	0.788 +/- 0.021	**0.629 +/- 0.015**

Cross-validation confirms the single test set findings. The difference between Logistic Regression and XGBoost is negligible (ΔAUC = 0.002), while Random Forest is substantially inferior (ΔAUC = 0.027-0.029). Notably, XGBoost achieves lower standard deviation on Recall (0.021 vs 0.035), indicating greater stability across folds.

6.4 McNemar's Test — Formal Model Comparison

McNemar's test was applied to formally compare Logistic Regression and Random Forest on the test set. Results: chi-squared = 9.301, p = 0.0023. The null hypothesis (that the two models make the same errors) is rejected at alpha = 0.001. However, this result reflects Random Forest's superior accuracy on the majority class (No Churn), not on the critical minority class (Churn). Recall remains the appropriate primary metric for this business problem.

7. Feature Importance & SHAP Analysis

7.1 Cross-Method Comparison

Feature importance was assessed using three methodologically independent approaches: SHAP values (game-theoretic, model-agnostic), Mean Decrease in Impurity (MDI, tree-based), and absolute logistic regression coefficients. Agreement across all three methods provides strong evidence for the robustness of the identified predictors.

Rank	SHAP (RF)	MDI (RF)	LR |Coefficient|
1	**tenure (0.110)**	**tenure (0.238)**	**Contract_Two year (1.398)**
2	InternetService_Fiber optic (0.062)	MonthlyCharges (0.217)	InternetService_Fiber optic (1.280)
3	Contract_Two year (0.061)	Contract_Two year (0.072)	tenure (0.758)
4	MonthlyCharges (0.060)	InternetService_Fiber optic (0.046)	Contract_One year (0.720)
5	Contract_One year (0.035)	PaymentMethod_Electronic check (0.038)	MonthlyCharges (0.463)

Four features appear in the top 5 across all three methods: tenure, InternetService_Fiber optic, Contract_Two year, and MonthlyCharges. This cross-method agreement constitutes strong evidence that these are genuine churn predictors, not model-specific artifacts.

7.2 SHAP Directional Findings

The SHAP beeswarm plot provides both magnitude and direction for each feature. Key findings:

Feature	Direction	SHAP Interpretation
tenure (low)	Increases churn	Blue dots right: low tenure strongly increases churn probability
tenure (high)	Reduces churn	Red dots left: high tenure strongly reduces churn probability
Fiber optic = Yes	Increases churn	Red dots right: having fiber optic increases churn risk (SHAP > 0)
Contract_Two year = Yes	Reduces churn	Red dots left: having 2-year contract strongly reduces churn risk
MonthlyCharges (high)	Increases churn	Red dots right: higher charges associated with increased churn risk

8. Threshold Optimization

The default classification threshold of 0.50 was evaluated against an optimal threshold derived from maximising the F1-score on the Precision-Recall curve. Results for the Logistic Regression model:

Threshold	Precision	Recall	F1-Score
0.50 (default)	0.511	0.778	0.617
0.596 (optimal F1)	0.548 (+0.037)	0.711 (-0.067)	0.619 (+0.002)

The optimal threshold improves Precision by +0.037 and F1 by +0.002, at the cost of Recall (-0.067). This trade-off translates to 60 fewer false alarms but 25 additional missed churners per 1,409 test cases. Given the asymmetric cost structure (the cost of a missed churner substantially exceeds the cost of a misdirected retention offer), the default threshold of 0.50 is recommended for production deployment in this context.

9. Business Recommendations

The following five recommendations are derived directly from the statistical and predictive modelling results. Each is supported by specific quantitative evidence from the analysis.

Recommendation 1: Contract Conversion Program

Evidence	Month-to-month churn rate: 42.7% vs 2.8% for two-year contracts. LR coefficient: -1.524 (OR = 0.218, $p < 0.001$). SHAP rank: #3.
Action	Offer targeted discounts (15-20%) or free service upgrades to month-to-month customers who commit to one- or two-year contracts. Prioritise customers with tenure < 12 months.
Impact	Converting 10% of month-to-month customers to two-year contracts could reduce overall churn rate by approximately 3-4 percentage points.

Recommendation 2: Early Lifecycle Intervention (0-12 Months)

Evidence	Median churned tenure: 10 months. SHAP rank: #1 (mean |SHAP| = 0.110). Mann-Whitney $p < 0.001$.
Action	Implement a structured 90-day onboarding programme with proactive check-ins and dedicated support. Deploy the churn model monthly to score all customers with tenure < 12 months.
Impact	Reducing early churn by 20% would save approximately 100 customers per month based on current churn patterns.

Recommendation 3: Fiber Optic Value Enhancement

Evidence	Fiber optic OR = 1.753 ($p = 0.005$). SHAP rank: #2 (mean |SHAP| = 0.062). MDI rank: #4.
Action	Conduct NPS surveys with fiber optic churners to identify pain points. Offer complimentary TechSupport (OR = 0.721) or OnlineSecurity (OR = 0.692) add-ons for 3 months to at-risk fiber customers.
Impact	Fiber optic customers generate higher revenue ($79+ per month average). Retaining even 15% of at-risk fiber customers would yield significant lifetime value recovery.

Recommendation 4: Payment Method Migration to Auto-Pay

Evidence	Electronic check OR = 1.548 ($p < 0.001$). SHAP rank: #6. MDI rank: #5.
Action	Offer a monthly discount (e.g., $5/month) for switching from electronic check to automatic bank transfer or credit card. Target specifically month-to-month customers with electronic check payment.
Impact	Auto-pay reduces active cancellation decisions and improves cash flow predictability, creating a dual retention and operational benefit.

Recommendation 5: Production Churn Scoring Deployment

Evidence	LR AUC = 0.845 +/- 0.014 (10-fold CV). Recall = 0.792 — identifies 79% of churners before they leave.
Action	Score all active customers monthly. Tier into: High Risk (P > 0.60) — direct account manager outreach; Medium Risk (0.30-0.60) — automated personalised offers; Low Risk (< 0.30) — standard loyalty communications.

Impact	Systematic risk-tiered intervention enables efficient allocation of the retention budget and is expected to reduce overall churn by 5-8 percentage points within 12 months.

10. Limitations

- **Cross-sectional snapshot:** Cross-sectional data: The dataset represents a single point in time and does not capture temporal dynamics, seasonal patterns, or competitor activity.
- **Causality:** No causal inference: Statistical associations do not imply causation. Observed relationships (e.g., fiber optic and churn) may be driven by unobserved confounders such as geographic market competition or price sensitivity segments.
- **Omitted variables:** Omitted variables: Key predictors such as service outage frequency, customer service call volume, NPS scores, and competitor pricing were not available. Their absence may introduce omitted variable bias in the coefficient estimates.
- **External validity:** Single split validation: Although 10-fold cross-validation was used, all results are based on a single dataset. External validation on a holdout period or a different market segment would strengthen generalisability.
- **RF interpretability:** Random Forest interpretability: MDI importance is known to be biased toward high-cardinality features. SHAP values were used to address this limitation, though SHAP TreeExplainer involves approximations for ensemble models.
- **Threshold optimality:** Threshold selection: The optimal threshold was selected based on F1-score maximisation, which assigns equal cost to false positives and false negatives. A formal cost-benefit analysis incorporating actual customer lifetime values would provide a more business-optimal threshold.

11. Conclusion

This report demonstrates that customer churn in this telecom dataset is predictable, statistically explainable, and actionable. Five analytical methods — visual EDA, formal statistical testing (Mann-Whitney U, Chi-square), Logistic Regression with confidence intervals, Random Forest with SHAP values, and XGBoost with cross-validation — converge on the same four primary churn drivers: contract type, customer tenure, internet service type, and monthly charges.

The best-performing models are Logistic Regression (AUC = 0.845 +/- 0.014, Recall = 0.792) and XGBoost (AUC = 0.847 +/- 0.021, Recall = 0.788), which are statistically equivalent in discriminative performance. Logistic Regression is preferred for business deployment due to its interpretability, coefficient significance, and confidence intervals. Random Forest, despite achieving the highest accuracy (0.781), is unsuitable as a churn detection tool due to its critically low Recall (0.477).

The five recommendations presented in Section 9 — contract conversion, early lifecycle intervention, fiber optic value enhancement, payment method migration, and production model deployment — constitute a data-driven retention strategy grounded directly in the quantitative evidence of this analysis. Implementation is estimated to reduce overall churn rate by 5-8 percentage points within a 12-month horizon.

End of Report